



Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE
100%

Module 3 Quiz

LATEST SUBMISSION GRADE

100%

1. A supervised learning model has been built to predict whether someone is infected with a new strain of a virus. The probability of any one person having the virus is 1%. Using accuracy as a metric, what would be a good choice for a baseline accuracy score that the new model would want to outperform?

1 / 1 point



Correct

2. Given the following confusion matrix:

1 / 1 point

	Predicted Positive	Predicted Negative
Condition Positive	96	4
Condition Negative	8	19

Compute the accuracy to three decimal places.



Correct

3. Given the following confusion matrix:

1 / 1 point

	Predicted Positive	Predicted Negative
Condition Positive	96	4
Condition Negative	8	19

Compute the precision to three decimal places.



Correct

4. Given the following confusion matrix:

1 / 1 point

	Predicted Positive	Predicted Negative
Condition Positive	96	4
Condition Negative	8	19

Compute the recall to three decimal places.



Correct

5. Using the fitted model `m` create a precision-recall curve to answer the following question:

1 / 1 point

For the fitted model `m`, approximately what precision can we expect for a recall of 0.8?

(Use `y_test` and `X_test` to compute the precision-recall curve. If you wish to view a plot, you can use `plt.show()`)

```
1 print(m)
```

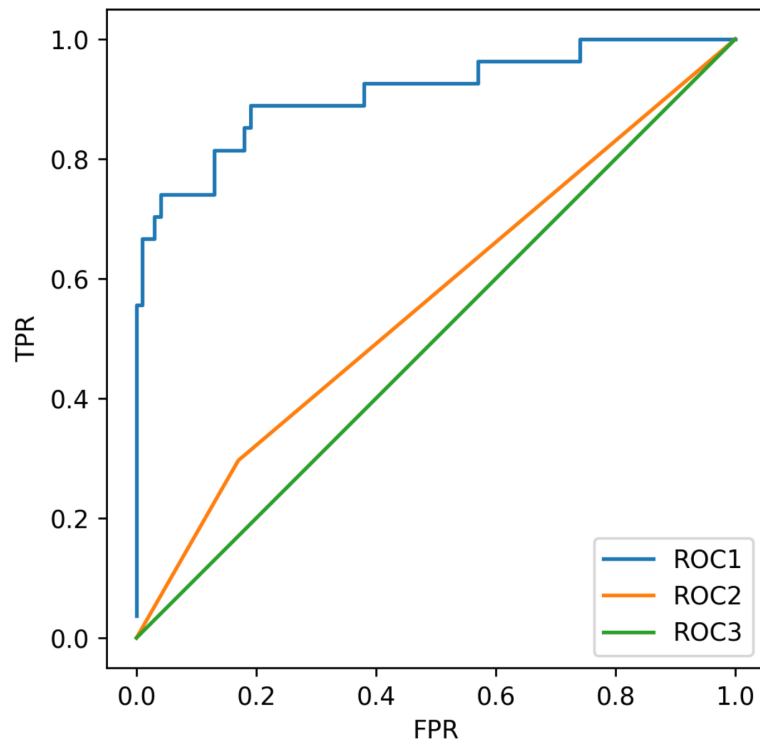
Run

✓ Correct

6. Given the following models and AUC scores, match each model to its corresponding ROC curve.

1 / 1 point

- Model 1 test set AUC score: 0.91
- Model 2 test set AUC score: 0.50
- Model 3 test set AUC score: 0.56

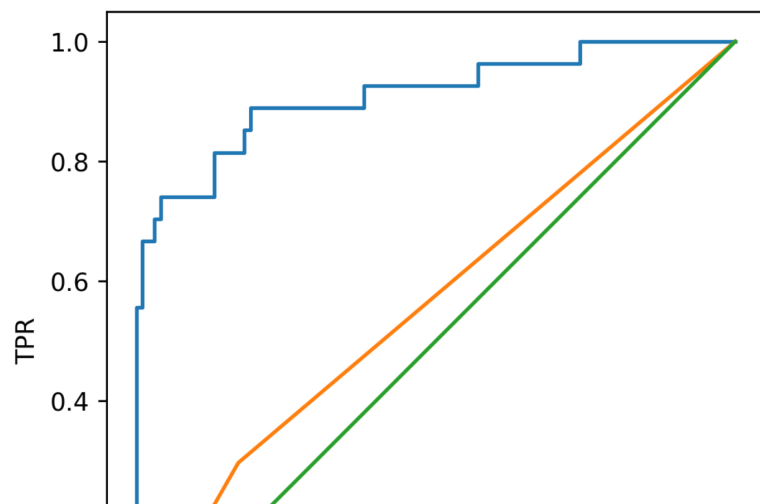


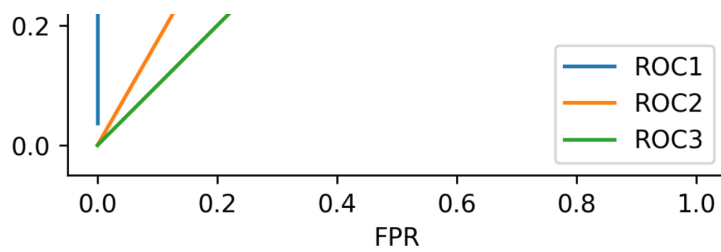
✓ Correct

7. Given the following models and accuracy scores, match each model to its corresponding ROC curve.

1 / 1 point

- Model 1 test set accuracy: 0.91
- Model 2 test set accuracy: 0.79
- Model 3 test set accuracy: 0.72





✓ Correct

8. Using the fitted model `m` what is the micro precision score?

1 / 1 point

(Use `y_test` and `X_test` to compute the precision score.)

```
1 print(m)
2 svm = m.fit(X_train, y_train)
3 svm_predicted_mc = svm.predict(X_test)
4 print(precision_score(y_test, svm_predicted_mc, average = 'micro'))
5
```

Run

Reset

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
0.744
0.744
None
```

✓ Correct

9. Which of the following is true of the R-Squared metric? (Select all that apply)

1 / 1 point

✓ Correct

10. In a future society, a machine is used to predict a crime before it occurs. If you were responsible for tuning this machine, what evaluation metric would you want to maximize to ensure no innocent people (people not about to commit a crime) are imprisoned (where crime is the positive label)?

1 / 1 point

✓ Correct

11. Consider the machine from the previous question. If you were responsible for tuning this machine, what evaluation metric would you want to maximize to ensure all criminals (people about to commit a crime) are imprisoned (where crime is the positive label)?

1 / 1 point

✓ Correct

12. A classifier is trained on an imbalanced multiclass dataset. After looking at the model's precision scores, you find that the micro averaging is much smaller than the macro averaging score. Which of the following is most likely happening?

1 / 1 point

✓ Correct

13. Using the already defined RBF SVC model `m`, run a grid search on the parameters `C` and `gamma`, for values `[0.01, 0.1, 1, 10]`. The grid search should find the model that best optimizes for recall. How much better is the recall of this model than the precision? (Compute recall - precision to 3 decimal places)

1 / 1 point

(Use `y_test` and `X_test` to compute precision and recall.)

```
1 print(m)
```

Run

Reset

✓ Correct

14. Using the already defined RBF SVC model `m`, run a grid search on the parameters C and gamma, for values [0.01, 0.1, 1, 10]. The grid search should find the model that best optimizes for precision. How much better is the precision of this model than the recall? (Compute precision - recall to 3 decimal places)

1 / 1 point

(Use y_test and X_test to compute precision and recall.)

```
1 print(m)
```

Run

Reset

✓ Correct